

● ENGAGEMENT PROPOSAL · OPERATING PLAN · PRE-ENGAGEMENT DRAFT

Teach the org to ship with AI& prove it moved the needle.

A two-dimensional engagement, in two coupled deliverables. The first upskills your engineers on AI-powered development in 25 hands-on hours. The second is the observability engine that makes the productivity gain legible to leadership in numbers, every day after.

01 · PROBLEM

The pressure is *specific* — and so is the gap.

Every engineering leader at a services firm is currently being squeezed on the same question: where are the AI productivity gains, and can you prove them? Clients won't accept "we're investing in it." Boards won't accept anecdote. Competitors who can quantify a delivery-velocity step-change have an asymmetric story in every procurement conversation.

The market response so far has mostly been one-off internal trainings — half-day workshops, recorded webinars, a few licenses for engineers who already liked tinkering. These leave no measurable trail and decay fast. What's missing isn't more training; it's training that's tightly coupled to the dashboard that will be used to judge whether it worked.

This proposal is the engagement just shipped at **ToTheNew**, generalized for [Client]. The shape is the same; the content adapts to your stack, cohort, and client mix.

02 · APPROACH

Two dimensions, *deliberately* coupled.

Training without measurement decays. Measurement without training shows nothing. The engagement runs both halves on a single timeline, so the metrics that prove the training worked are the same metrics the team learns to influence.

DIMENSION 1

AI Learning Curriculum

10 days × 2.5 hours, hands-on, Claude Code + Cursor.
Zero-to-hero arc designed for engineers who already use Copilot casually and need to become productive with agentic workflows on real brownfield codebases.

CONSUMER · ENGINEERING COHORT

DIMENSION 2

Observability & Adoption Engine

24×7 dashboard against four categories — Activity, Throughput, Adoption, Quality — that defends the training to leadership. Every metric is grounded in industry practice (DORA, GitHub/Anthropic enterprise admin telemetry).

CONSUMER · ENGINEERING LEADERSHIP

03 · DIMENSION 1

Curriculum spine — *ten* days to working AI-native engineers.

One shared spine, hands-on heavy, role-forked only where it must be (Day 8). Each day = 30 min framing → 90 min lab → 30 min debrief. The arc compounds — later days depend on earlier ones landing.

DAY	THEME	WHAT LANDS	TOOL
01	Mental models & setup	Both tools running; first prompts that work; the shift from chat-paste to in-IDE.	BOTH
02	Context engineering ★ biggest unlock	"Context first, prompt second" — file tagging, CLAUDE.md, .cursorrules in muscle memory.	BOTH
03	Inline + chat workflows	Cursor fluency — Tab, ⌘K, chat, composer; right tool for each task.	CURSOR
04	Agentic workflow	First real agentic delivery; mental shift from "AI completes my code" to "AI executes, I supervise."	CLAUDE CODE
05	Planning before coding	A repeatable pre-flight discipline; the plan becomes the PR description.	CLAUDE CODE
06	Brownfield: read, modify	Comfort in unfamiliar code with AI as guide; catching hallucinated function calls.	BOTH
07	Tests, debugging, review	AI at quality surfaces without quality going down; spotting AI-code smells in five seconds.	BOTH
08	Stack-forked labs	Universal skills meet role patterns: FE component, BE endpoint, DevOps IaC, Data SQL/DAG.	BOTH
09	Customization & guardrails ◆ dim-1 meets dim-2	Committed CLAUDE.md / .cursorrules / custom commands in real repos. Durability.	BOTH
10	Capstone & integration	Shipped PR per attendee; team norms for AI-generated code; one habit each commits to.	BOTH

04 .
DIMENSION 2

Observability scaffold — *four* categories, grounded in industry practice.

24 metrics across four categories, each instrumented from an established source (Git host, JIRA, IDE telemetry, Anthropic billing). The scaffold is intentionally redundant — every velocity number pairs with a quality counterweight, so leadership can't be sold speed at the cost of stability.

Activity

The raw volume signals. Easy to game alone; valuable as denominators for the categories below.

- **LOC / dev / year** — the perf-review datapoint at Google, Meta, Amazon
- **PRs raised, CLs reviewed, design docs reviewed**
- **Time spent coding** · agent vs user hours

Adoption

Are tools actually being used, with the configurations that make them productive? Decay shows here first.

- **Cursor DAU, Claude Code DAU**
- **CLAUDE.md coverage** — leading indicator of durable productivity
- **Token consumption, custom commands, MCP server usage**

Throughput

"Did we get faster" — every line here pairs with a quality line. The DORA metrics anchor this category.

- **Ticket cycle time** · DORA p50 — the headline metric
- **SWE hours saved, features shipped per sprint**
- **Brownfield velocity, plan-first ratio**

Quality

The counterweight to throughput. If speed celebrates while quality silently drops, the engagement did damage, not value.

- **PR rework rate, post-merge bug rate** · DORA change-failure
- **AI-touch ratio on bugs** · the "is AI making us worse?" answer
- **Test coverage on AI code, hallucination incidents, context-grounded prompt ratio**

05 · MODEL

Engagement model — *timeline*, deliverables, your part.

TIMELINE	DELIVERABLES	WHAT [CLIENT] PROVIDES
<i>Week 0 · Baseline</i> Pre-engagement survey, license confirmation, repo access, two weeks of pre-training measurement so post-training has something to compare against.	— Cohort onboarding survey + baseline report — 10 days of facilitated training (in-person) — Per-attendee CLAUDE.md + .cursorsrules committed to at least one of your real repos	— Cohort of engineers, blocked on calendars for the 10 days — Cursor + Claude Code licenses for the cohort — during AND for 90+ days after — Access to representative codebases for Days 6, 8, 9
<i>Weeks 1–2 · Delivery</i> The 10-day curriculum, offline, 2.5 hrs/day. Real TTN-style tickets sourced from your codebase for Days 6, 8, 10.	— Per-attendee shipped PR (Day 10 capstone) — Observability dashboard wired against your Git host, JIRA, IDE telemetry, Anthropic billing	— Real tickets from your backlog for Days 6, 8, 10 — Engineering sponsor empowered to commit CLAUDE.md to those repos
<i>Weeks 3+ · Sustain</i> Dashboard goes live. Bi-weekly metric review for the first quarter. Curriculum artifacts (CLAUDE.md, custom commands) maintained in your repos by your team.	— Metric documentation + 1-pager for leadership — 90-day post-engagement check-in	— API access to Git host, JIRA, IDE admin for instrumentation — Two-week pre-engagement window for baseline measurement

06 · OPEN

Open questions before we lock the plan.

Decisions that, once answered, move this from a proposal into an executable engagement. None block design; all block delivery.

- 01

Cohort size, composition, and split.

One shared track works to ~30 engineers per cohort. Above that, parallel sessions. Stack mix decides whether Day 8 forks four ways or fewer.
- 02

License coverage during and after.

If engineers fall back to free tiers after Week 2, the dashboard numbers will flatten and the engagement will look like it failed.

03

Real-problem sourcing.

One unfamiliar repo for Day 6, four small role-specific tasks for Day 8, one real ticket per attendee for Day 10. Toy problems lose the room.

04

Baseline measurement window.

Two weeks of pre-training instrumentation so the after-picture has something to compare against. Earlier the better.

05

Metric mix calibration.

What [Client] already tracks vs. what we'd add. The dashboard inherits your existing trust where possible; nothing duplicated.

"The training has to teach the right skills, and the dashboard has to measure the right things — or neither dimension does its job."

OUT-001 • REV 0.1

PREPARED BY [SENDER]

REFERENCE ENGAGEMENT • TOTHE NEW